

FEBRUARY 2026

THE CLAUDE BEAST

Complete Overview: 17 Topics for
Mastering Agentic Development

Optimized for start.io Agentic Hackathons

System Prompt: CLAUDE.md & AGENTS.md

Episode 1 | The Foundation of Every Session

- **CLAUDE.md = persistent project memory** — loaded into Claude's system prompt every session
- **6-level memory hierarchy** — managed policy, user, project, rules, local, auto memory
- **AGENTS.md = cross-tool open standard** — works with Cursor, Copilot, Zed, and 10+ tools
- **Symlink pattern** — enables interoperability between CLAUDE.md and AGENTS.md
- **Best practice: keep it concise** — under 300 lines, progressive disclosure via @imports
- **/init command** — auto-generates a starter CLAUDE.md from your codebase

Key Insight

Without CLAUDE.md, every session starts from zero. It is the only file loaded by default into every conversation with the agent.

PRD: Product Requirements Documents

Episode 2 | From Human Docs to AI Programming Interfaces

- **PRDs evolved** — from human-readable docs to AI programming interfaces
- **Claude Code uses 3 context sources** — CLAUDE.md + codebase + conversation history
- **Gold standard methodology** — PRD → Plan → Todo for structured development
- **8+ skills and plugins available** — for automated PRD generation
- **PRDs solve context overload** — give Claude a complete spec before coding begins
- **Bridges the gap** — between business requirements and agentic execution

Key Insight

The PRD is not just documentation. It is the programming interface that tells Claude what to build, how, and why.

Claude Skills

Episode 3 / Markdown Workflow Templates with Progressive Disclosure

- **Markdown-based workflow templates** — with progressive disclosure architecture
- **Open standard** — adopted by 26+ platforms (agentskills.io registry)
- **Key skills** — Superpowers (60K stars), frontend-design, Vercel React best practices
- **Enforce structured workflows** — brainstorm → plan → TDD → review cycle
- **Install with one command** — /install-skill or direct GitHub URL
- **Skills complement CLAUDE.md** — skills are task workflows, CLAUDE.md is project context

Key Insight

Skills are reusable expert templates that enforce quality workflows. They turn Claude from a general assistant into a specialized practitioner.

Claude Hooks

Episode 4 | Deterministic Control Over Probabilistic AI

- **14+ lifecycle events** — across Session, Tool, Agent, and User categories
- **3 hook types** — command (shell), prompt (LLM injection), agent (multi-turn)
- **CLAUDE.md = guidance (suggestions)** — vs Hooks = enforcement (guaranteed execution)
- **Pre/post patterns** — validate before actions, transform after results
- **Real-world uses** — auto-format on save, block dangerous commands, inject context
- **Configured in .claude/settings.json** — with matcher patterns for targeted activation

Key Insight

CLAUDE.md suggests behavior. Hooks enforce it. They provide deterministic control over Claude's probabilistic actions.

LSP: Language Server Protocol

Episode 5 | IDE-Level Code Understanding for Claude Code

- **9 LSP operations** — hover, goToDefinition, findReferences, diagnostics, and more
- **Performance leap** — 900x faster navigation, 75% token reduction in prompts
- **Multi-language support** — 11+ official languages, 17+ community languages
- **Type-aware intelligence** — Claude sees types, signatures, and cross-file references
- **Automatic activation** — Claude Code starts LSP servers for detected languages
- **Reduces hallucination** — Claude navigates code precisely instead of guessing

Key Insight

LSP gives Claude the same code intelligence that IDEs give human developers: jump to definition, find references, and see type information.

Claude Plan Mode

Episode 6 | Research Before You Code

- **Read-only exploration mode** — research and plan without modifying any files
- **Activation** — Shift+Tab twice, /plan command, or --permission-mode plan flag
- **opusplan alias** — Opus 4.6 for planning, Sonnet for execution (cost optimization)
- **Integrates with Agent Teams** — plan mode required for controlled team members
- **Pairs with Superpowers** — plan phase naturally feeds into TDD execution
- **Prevents premature coding** — forces analysis before implementation

Key Insight

Plan Mode forces a "think before you code" discipline. Use Opus 4.6 to plan, Sonnet to execute.

SuperPowers, GSD, OpenSpec & Spec Kit

Episode 7 | Four Tools Ending the "Vibe Coding" Era

- **Superpowers (59.8K stars)** — enforced TDD + brainstorm-plan-review workflow
- **GSD (19.7K stars)** — meta-prompting + context engineering for deep work
- **OpenSpec (25.4K stars)** — spec-driven development with structured templates
- **Spec Kit (71.4K stars)** — specification management and validation framework
- **Common philosophy** — structure first, code second, review always
- **All installable as Claude Skills** — one command setup

Key Insight

These four tools share one philosophy: stop vibe coding. Structure your intent before writing a single line of code.

EXA MCP: AI-Native Search

Episode 8 | Real-Time Web Intelligence via MCP Protocol

- **AI-native search engine** — \$700M valuation, purpose-built for LLM consumption
- **MCP = "USB-C for AI"** — universal tool protocol created by Anthropic
- **8+ tools** — web search (2M+ calls), code context, deep research agent
- **One-command setup** — free tier includes \$10 in API credits
- **Deep Researcher** — autonomous multi-step research with synthesis
- **Code Context tool** — finds documentation, examples, and API references

Key Insight

MCP is the universal protocol for connecting AI to external tools. EXA brings live web knowledge into Claude's context.

Atlassian MCP: Jira, Confluence & Compass

Episode 9 | Enterprise Project Management for Claude Code

- **Bridges Claude Code** — to Jira, Confluence, and Compass ecosystems
- **35+ tools** — 13 Jira + 11 Confluence + 6 Compass + platform tools
- **Official (Rovo) vs Community** — enterprise Rovo agent or sooperset (1.2K stars)
- **PRD from Confluence** — generate product requirements from existing wiki pages
- **Sprint planning** — create, estimate, and assign Jira issues from conversation
- **Context-aware coding** — Claude reads Jira tickets and Confluence docs while coding



Key Insight

Atlassian MCP closes the loop between project management and code. Claude reads your tickets and docs, then builds what they describe.

Claude Sub Agents

Episode 10 | Isolated AI Workers via the Task Tool

- **Isolated AI workers** — spawned via the Task tool with their own context window
- **Built-in types** — Explore (read-only), Plan (analysis), general-purpose (full access)
- **Context isolation** — each agent gets a fresh context window, preventing pollution
- **Worktree isolation** — git worktrees enable true parallel file editing without conflicts
- **Hub-and-spoke architecture** — results flow back to parent agent for synthesis
- **Custom agents** — define specialized .md files in .claude/agents/ directory

Key Insight

Sub agents give you parallelism with isolation. Each agent works in its own context window, preventing token pollution across tasks.

Claude Agent Teams

Episode 11 | Multi-Agent Collaboration with Peer-to-Peer Messaging

- **Multi-agent collaboration** — Team Lead + Teammates + Shared Tasks + Mailbox
- **7 coordination primitives** — create, message, broadcast, task, assign, shutdown, plan
- **Peer-to-peer messaging** — agents communicate directly, not just through the lead
- **Proven at scale** — 16 agents built a 100K-line C compiler (\$20K API cost)
- **Use cases** — parallel code review, competing hypotheses, cross-layer development
- **Shared task list** — all agents read/write the same task queue for coordination

Key Insight

Agent Teams are Sub Agents with communication. They can message each other, share tasks, and collaborate on complex problems.

Greptile: AI Code Review

Episode 12 | Full Codebase Context with Graph-Based Analysis

- **AI code review** — with full codebase context using graph-based indexing
- **v3 powered by Claude Opus 4.5** — 256% improvement over previous versions
- **82% bug catch rate** — 100% detection on high-severity bugs
- **Autonomous greploop** — review → fix → push → re-review until clean
- **MCP integration** — use Greptile tools directly from Claude Code
- **Custom context rules** — teach Greptile your team's patterns and standards

Key Insight

Greptile does not just lint individual files. It understands your entire codebase graph and catches architectural issues.

/team-review: Agent Team Code Review

Episode 13 | Three Parallel AI Reviewers, One Comprehensive Report

- **Three specialist reviewers** — security, performance, and test coverage in parallel
- **Built on Agent Teams** — each reviewer is an independent Claude instance
- **Zero configuration** — just type /team-review, auto-detects branch changes
- **Synthesized report** — findings grouped by severity with remediation suggestions
- **Part of startappdev/agents** — MIT licensed, pairs with /greptile-review-loop
- **Multi-pass philosophy** — specialist depth beats single-pass breadth

Key Insight

/team-review deploys three focused reviewers simultaneously. Each has its own context window, preventing cross-contamination between dimensions.

/greptile-review-loop

Episode 14 | Autonomous Review-Fix-Merge for Every Pull Request

- **Fully autonomous closed loop** — review → fix → push → re-review → merge
- **Parallel code-fixer agents** — one per affected file, running simultaneously
- **Five hard stop conditions** — score thresholds, iteration caps, build validation
- **HANDLED_IDS tracking** — prevents re-fixing previously addressed comments
- **Auto-merge at score $\geq 4/5$** — with squash-merge and branch cleanup
- **Safety by design** — never merges below 3/5, never loops beyond 5 iterations

Key Insight

One command, zero human intervention. PR is created, reviewed, fixed, and merged automatically until quality thresholds are met.

Claude Worktrees

Episode 15 | Parallel Isolated Development Sessions via Git Worktrees

- **One-command isolation** — `claude --worktree feature-name` creates isolated workspace
- **Built on git worktree** — shared `.git` history, separate files and branches
- **Auto-cleanup** — worktree and branch removed if no changes on exit
- **Subagent integration** — `set isolation: worktree` in agent frontmatter
- **Foundation for parallelism** — enables safe file editing across sub-agents and teams
- **Cross-platform** — CLI, Desktop, VS Code, JetBrains, web, and mobile

Key Insight

Worktrees are the file-isolation layer that makes all other parallelism safe. Both sub-agents and agent teams use worktrees under the hood.

Portless: Named URLs for Local Dev

Episode 16 | Stable .localhost URLs for Humans and AI Agents

- **Named .localhost URLs** — `myapp.localhost:1355` instead of `localhost:3000`
- **Proxy-based routing** — single entry point on port 1355, random backend ports
- **11 frameworks tested** — Next.js, Vite, Nuxt, Astro, Angular, and more
- **HTTPS / HTTP2** — auto-generated TLS certificates with local CA
- **AI agent ready** — stable semantic endpoints agents can reliably target
- **2,400+ stars in 10 days** — from Vercel Labs, Apache-2.0 licensed

Key Insight

As AI coding agents become standard, they need stable semantic URLs for local services. Portless gives them -- and humans -- exactly that.

Handy: Offline Speech-to-Text

Episode 17 | Free, Private, Local Transcription for Developers

- **100% offline speech-to-text** — no audio ever leaves your machine
- **Multiple models** — Whisper (GPU) and Parakeet V3 (CPU, ~5x real-time)
- **Voice-to-AI workflow** — speak prompts to Claude Code across terminals
- **LLM post-processing** — optional AI cleanup before text insertion
- **16K+ stars, 900K downloads** — featured in WIRED, MIT licensed
- **Tauri + Rust** — native performance, ~12.6 MB installer, not Electron

Key Insight

Speaking is 3-4x faster than typing. Handy turns voice into the natural interface for agentic development with Claude Code.

Go Build Something Amazing



1. System Prompt (CLAUDE.md)
2. PRD & Requirements
3. Claude Skills
4. Claude Hooks
5. LSP Integration
6. Plan Mode
7. SuperPowers / GSD / Spec
8. EXA MCP Search
9. Atlassian MCP
10. Sub Agents
11. Agent Teams
12. Greptile Code Review
13. /team-review
14. /greptile-review-loop
15. Claude Worktrees
16. Portless
17. Handy STT