

THE CLAUDE BEAST

Greptile: AI Code Review

Full codebase context. Agentic architecture. Autonomous quality loops.

2,000+ organizations | 1B+ lines reviewed | 1M+ issues caught/month

What Is Greptile?

An AI code review agent that understands your entire codebase -- not just the diff

What It Catches

- Logic errors: bugs, edge cases, race conditions
- Security vulnerabilities: SQL injection, unvalidated inputs
- Performance issues: inefficient patterns, impact analysis
- Style violations: naming, dead code, complexity
- Code repetition across the entire repository

Key Differentiators

- Full codebase graph -- every function, class, dependency
- Conversational: reply with @greptileai for follow-ups
- Reinforcement learning from team feedback
- Cross-repo context via pattern repos
- Custom rules in plain English

Anatomy of a Greptile Review

Review Process (~3 min)

- PR detected → Greptile begins analysis
- Builds full codebase context (not just the diff)
- Posts PR summary + inline comments
- Auto-generated Mermaid diagrams per change type

PR Summary Includes

- Plain-language explanation of what the PR does
- File-by-file breakdown with issue counts
- Inline comments with click-to-apply code suggestions

Confidence Score (0-5)

5/5

Production ready

→ Merge

4/5

Minor polish needed

→ Small fixes, then merge

3/5

Implementation issues

→ Address feedback first

2/5

Significant bugs

→ Needs rework

0-1

Critical problems

→ Major rethink

Claude Code MCP Plugin: 11 Tools

Fetch reviews, apply fixes, search feedback -- all without leaving the terminal

Pull Request Tools (3)

```
list_pull_requests  
get_merge_request  
list_merge_request_comments
```

Code Review Tools (3)

```
trigger_code_review  
get_code_review  
list_code_reviews
```

Comment Search (1)

```
search_greptile_comments  
(across all repos,  
by query,  
date, addressed  
status)
```

Custom Context (4)

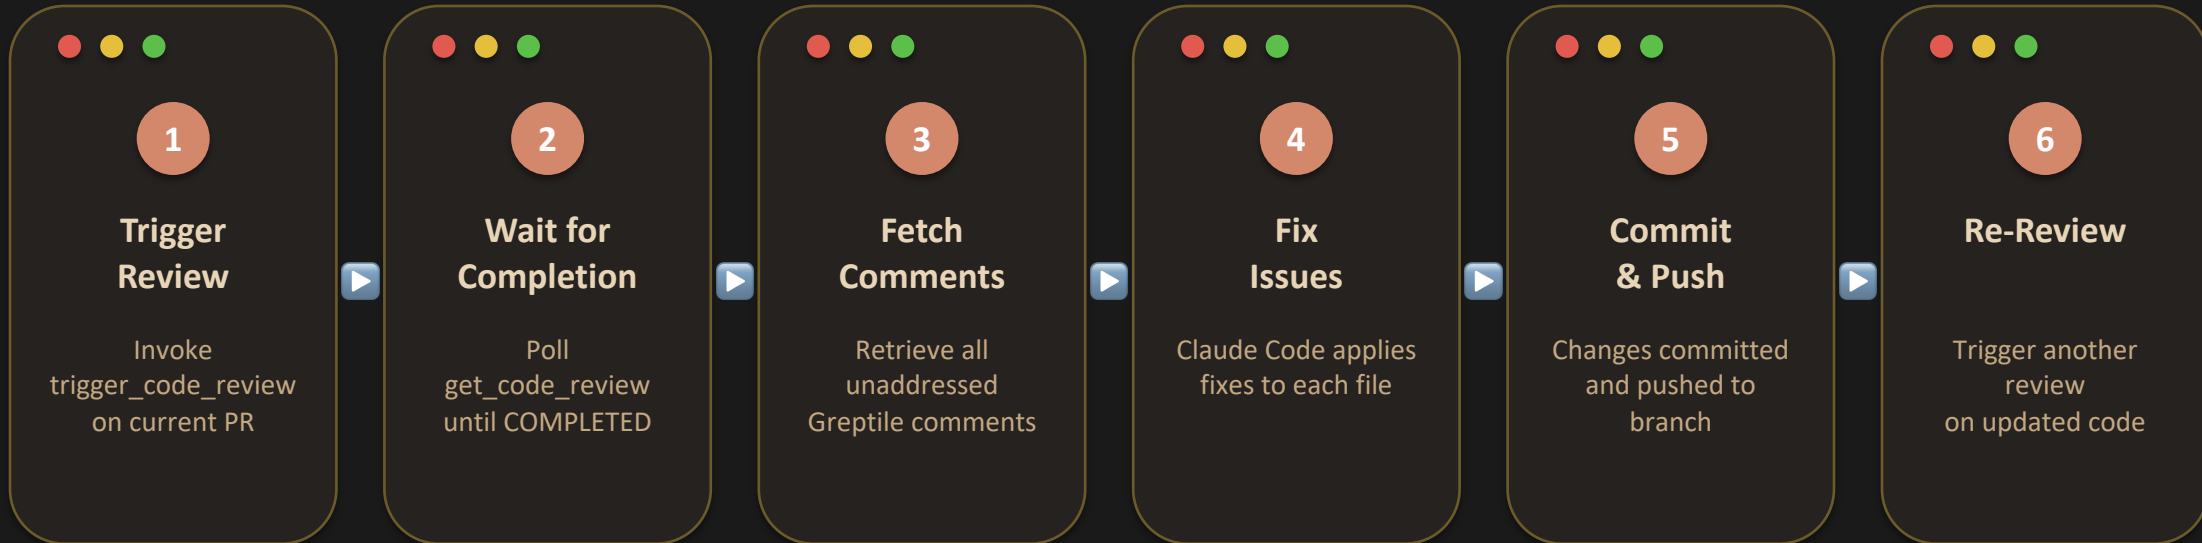
```
create_custom_context  
search / get / list  
custom_context
```

Example Prompts

- “Show me all Greptile comments on my open PR and fix them”
- “Run a Greptile review on PR #8”
- “What are the most common issues Greptile has flagged this month?”

The Greploop: Autonomous Review-Fix-Push

Loop until 5/5 confidence and zero unaddressed comments



↺ Repeats until: confidence = 5/5 AND unaddressed comments = 0

The Vision: Human expresses intent → Coding agent executes → Review agent validates → Loop until approved

Learning from Your Team's Feedback



1

Reading Team Comments

Observes patterns in your team's PR discussions and extracts preferences



2

Thumbs Up/Down

Direct feedback on which comments are valuable vs. noise



3

Custom Rules

Explicit standards defined in plain English by your team



4

Pattern Repos

Cross-repository context for related codebases



Nitpick Reduction

After 2-3 weeks, Greptile stops commenting on things your team doesn't care about. Signal over noise.

Why Greptile Matters

The Quality Gap Is Real

82% bug catch rate, 100% on high-severity. 41% ahead of the next best tool. Code review at scale demands AI.

Agentic Architecture Changes Everything

v3's detective approach delivers +256% improvement. Claude Opus 4.5 + Agent SDK = deep, context-aware reviews.

The Autonomous Loop Is the Future

Claude Code + Greptile = closed-loop quality. /greploop runs until 5/5 confidence. Zero context-switching.

"The Agent SDK has allowed us to ship faster with far more cost effectiveness." -- Daksh Gupta, CEO, Greptile