

THE CLAUDE BEAST

Handy: Offline Speech-to-Text

Free, Private, Local Transcription for Developers and AI Agents

16K+ Stars | 900K Downloads | Featured in WIRED | MIT License

What Is Handy?

Press a shortcut. Speak. Text appears. All local.

Core Concept

A desktop application that converts speech to text using local AI models. Press a global keyboard shortcut, speak, and your words appear in whatever application is active -- all processed on-device with zero cloud dependency.

Key Features

- Global keyboard shortcuts (push-to-talk or toggle)
- Voice Activity Detection (Silero VAD v4)
- Multiple transcription models (Whisper + Parakeet)
- LLM post-processing for cleanup and formatting
- Replacement rules with import/export
- CLI remote control for scripting
- Cross-platform: macOS, Windows, Linux

Origin Story

CJ Pais built Handy after breaking a finger and needing one-handed computer operation. Existing STT tools were either proprietary, cloud-dependent, or not extensible.


Philosophy: "The Most Forkable Speech-to-Text App" -- not the best, but the most extensible and modifiable.

Now used by developers with dystonia and other mobility limitations worldwide.


Architecture: Tauri + Rust + React

Native performance, cross-platform, ~12.6 MB installer

Processing Pipeline



Core Libraries



GPU: Metal (macOS) | Vulkan (Windows/Linux) | OpenBLAS (Linux fallback) | Language: Rust 49% + TypeScript 49%

Transcription Models

Choose between GPU-accelerated accuracy or CPU-optimized speed

 Whisper Small 487 MB GPU Good accuracy, smaller footprint	 Whisper Medium 492 MB GPU Balanced accuracy and size	 Whisper Turbo 1.6 GB GPU Fast inference with GPU	 Whisper Large 1.1 GB GPU Highest accuracy (multilingual)	 Parakeet V3 CPU-opt CPU ~5x real-time on mid-range hardware	 Custom GGML Variable GPU/CPU Drop any .bin file into models dir
--	---	--	---	--	--



Community Favorite: Parakeet V3

Parakeet V3 is the standout model for developers. CPU-optimized (no GPU needed), automatic language detection, and near-instant transcription at ~5x real-time speed. Multiple Hacker News users switched from paid tools specifically because of Parakeet's speed.

Additional Language Support

- Breeze ASR and SenseVoice Small for Chinese language transcription
- Translate-to-English setting for automatic translation from any language
- N-gram matching for multi-word custom word correction

LLM Post-Processing & Customization

LLM Post-Processing (Debug Menu)

- Send transcripts through local or remote LLM
- Custom system prompts for cleanup rules
- Context variables: `$currentApp`, `$shortPrevTranscript`
- Toggle via keyboard shortcut (`--toggle-post-process`)
- Accessible via `Cmd/Ctrl+Shift+D` debug menu

Replacement Rules

- Pattern-based find-and-replace
- Import/export via JSON
- Pre-built templates: English, French punctuation
- N-gram matching for multi-word corrections
- Custom technical vocabulary handling

CLI & Scripting Integration

<code>--toggle-transcription</code>	Start/stop recording from scripts
<code>--toggle-post-process</code>	Enable/disable LLM post-processing
<code>--cancel</code>	Cancel current transcription
<code>--start-hidden --no-tray</code>	Autostart without visible UI
<code>SIGUSR1 / SIGUSR2</code>	Unix signals for window managers (Sway, i3, KDE)

Adoption: 16K Stars, 900K Downloads

One of the fastest-growing open-source desktop apps in the STT space



16,037

GitHub Stars



886K+

Total Downloads



80

Contributors



30

Releases

Press & Community

WIRED

"Stop Using Your Keyboard and Start Using This Simple, Free Speech-to-Text App"

Hacker News

247 points, 110 comments -- praised for simplicity, speed, and privacy

GIGAZINE

Featured in Japan's leading tech publication

Homebrew

Available via: `brew install --cask handy`

Sponsors: Wordcab | Epicenter | BoltAI | v0.7.8 latest (Feb 25, 2026) | 5 releases in Feb 2026

Handy vs. Alternatives

Handy

Free, open-source, offline
MIT license, 16K+ stars
Tauri + Rust (small footprint)
Multiple models, extensible

SuperWhisper

Paid (\$9.99+/month)
Mac only, closed source
Cloud + local options
Polished UI

VoiceInk

Paid (\$29 one-time)
Mac only, closed source
Local processing
Apple Silicon optimized

WisprFlow

Paid subscription
Multi-platform
Cloud-dependent
Integrations focused

Handy's Key Differentiators

- 100% offline: no audio ever leaves your machine. No cloud, no subscriptions, no internet required.
- Truly free and open source (MIT): no feature gating, no freemium limits
- AI-agent friendly: includes CLAUDE.md and AGENTS.md for Claude Code modification
- CLI-controllable: scriptable from window managers, automation tools, and hooks
- "The most forkable STT app": designed for modification, not vendor lock-in

Claude Code Integration

Agent-Native Codebase

- CLAUDE.md in repo root for Claude Code guidance
- AGENTS.md for cross-tool compatibility
- Full architecture documented for AI modification
- Code style and commit conventions specified
- Platform-specific notes for AI understanding

Integration Patterns

- Voice → Claude Code: speak prompts to terminal
- Voice → Multiple worktrees: fast context switching
- Voice → Agent teams: orchestrate via speech
- Post-process → Claude: clean up technical speech
- Hooks → Handy: trigger via Claude Code events

The Complete Developer Voice Stack

Handy (speech input) → Claude Code (execution) → Worktrees (isolation) → Agent Teams (coordination)

The developer speaks intent. The system translates to text, routes to the right terminal, and Claude Code executes. Voice becomes the natural interface for agentic development.

Why Handy Matters

Voice Is the Fastest Input

Speaking is 3-4x faster than typing. For complex multi-step AI prompts, voice input is a genuine productivity multiplier across multiple terminals and worktrees.

Privacy Without Compromise

100% offline, no cloud, no subscriptions. 900K downloads prove that developers will choose privacy when the quality is good enough -- and Parakeet V3 is very good.

The Agent-Native Future

Handy + Claude Code + Worktrees + Agent Teams = a voice-driven development environment. Speak your intent. Agents handle the rest.

github.com/cjpais/Handy | handy.computer | MIT License | `brew install --cask handy`